



# **Support Vector Machines in R: a benchmark study**

David Meyer

June 26, 2003

# Overview

---

1. Support Vector Machines
2. `libsvm`
3. The R interface
4. Benchmark results
5. The EUNITE competition

# SV $\epsilon$ -Regression

---

We are given training data:


$$\{(\mathbf{x}_i, y_i)\}_1^l, \mathbf{x}_i, y_i \in \mathbb{R}^n$$

Suppose the data can be explained by a linear model.

Goal: find a fitting hyperplane  $\langle \mathbf{w}, \mathbf{x} \rangle + b = 0$


# SV $\epsilon$ -Regression

---




# SV $\epsilon$ -Regression

---




# SV $\epsilon$ -Regression

---




# SV $\epsilon$ -Regression

---



# SV $\epsilon$ -Regression





# SV $\epsilon$ -Regression

---

The distance from the hyperplane to a border is:

$$\frac{1}{\|\mathbf{w}\|}$$

Thus, the margin is given by:

$$\frac{2}{\|\mathbf{w}\|}$$

## Goal:

Maximize margin such it will contain all points with errors smaller than an  $\epsilon$  to be specified.

# SV $\epsilon$ -Regression

---


Formally:

$$\min_{\mathbf{w}, b} \frac{\|\mathbf{w}\|^2}{2}$$

under constraints:

$$\begin{aligned} y_i - \langle \mathbf{w}, \mathbf{x}_i \rangle - b &\leq \epsilon \\ \langle \mathbf{w}, \mathbf{x}_i \rangle + b - y_i &\leq \epsilon \end{aligned}$$

# SV $\epsilon$ -Regression



# SV $\epsilon$ -Regression

---

## Handling outliers:

Allow points to be “outside” the margin by using weights  $\xi_i$ :

$$\begin{aligned} y_i - \langle \mathbf{w}, \mathbf{x}_i \rangle - b &\leq \epsilon + \xi_i \\ \langle \mathbf{w}, \mathbf{x}_i \rangle + b - y_i &\leq \epsilon + \xi_i^* \\ \xi_i, \xi_i^* &\geq 0, i = 1, \dots, l \end{aligned}$$

but introduce penalty weight  $C$ :

$$\min_{\mathbf{w}, b, \xi, \xi^*} \frac{\|\mathbf{w}\|^2}{2} + C \sum_{i=1}^l (\xi_i + \xi_i^*)$$

# SV $\epsilon$ -Regression

## Handling non-linearity:

Map the input data  $\mathbf{x}_i$  into a (possibly higher dimensional) *feature space*:

$$\mathbf{x}_i \mapsto \phi(\mathbf{x}_i)$$

Thus, the complete optimization problem becomes:

$$\min_{\mathbf{w}, b, \xi, \xi^*} \frac{\|\mathbf{w}\|^2}{2} + C \sum_{i=1}^l (\xi_i + \xi_i^*)$$

$$\begin{aligned} \text{s.t.} \quad y_i - \langle \mathbf{w}, \phi(\mathbf{x}_i) \rangle - b &\leq \epsilon + \xi_i \\ \langle \mathbf{w}, \phi(\mathbf{x}_i) \rangle + b - y_i &\leq \epsilon + \xi_i^* \\ \xi_i, \xi_i^* &\geq 0, i = 1, \dots, l \end{aligned}$$

# SV $\epsilon$ -Regression

## The Kernel trick:

Find solution using Lagrangian method in the *dual* form:

$$\begin{aligned}
 \min_{\alpha, \alpha^*} \quad & \frac{1}{2} \sum_{i, j=1}^l (\alpha_i - \alpha_i^*)(\alpha_j - \alpha_j^*) \langle \phi(\mathbf{x}_i), \phi(\mathbf{x}_j) \rangle \\
 & + \epsilon \sum_{i=1}^l (\alpha_i + \alpha_i^*) + \sum_{i=1}^l y_i (\alpha_i + \alpha_i^*) \\
 \text{s.t.} \quad & \sum_{i=1}^l (\alpha_i - \alpha_i^*) = 0 \\
 & 0 \leq \alpha_i, \alpha_i^* \leq C, \quad i = 1, \dots, l
 \end{aligned}$$

The inner product  $K(\mathbf{x}_i, \mathbf{x}_j) = \langle \phi(\mathbf{x}_i), \phi(\mathbf{x}_j) \rangle$  is called a *Kernel*.

# SV $\epsilon$ -Regression

---

## Kernels:

For certain functions  $\phi$ , the inner product can be computed by  $K(\mathbf{x}_i, \mathbf{x}_j)$  without explicitly knowing  $\phi(\cdot)$ .

Some popular kernels:

Polynomial	$(\gamma \langle \mathbf{x}_i, \mathbf{x}_j \rangle + \beta)^d$
Radial Basis Function	$\exp(-\gamma \ \mathbf{x}_i - \mathbf{x}_j\ ^2)$
Sigmoid	$\tanh(\gamma \langle \mathbf{x}_i, \mathbf{x}_j \rangle + \beta)$

We do have to specify (in addition to  $C$  and  $\epsilon$ ) the hyperparameters  $d$ ,  $\gamma$ , and  $\beta$ : this usually requires *tuning*.

# SV $\epsilon$ -Regression

---

## Predicting:

The approximate function is simply:

$$f(x) = \sum_{i=1}^l (-\alpha_i + \alpha_i^*) K(x_i, x) + b$$

... a weighted sum (i.e., superposition) of kernel functions.



# Chih-Chen Lin's libsvm

---

- ❄ Starting point:  $\text{svm}^{\text{light}}$  (Joachims, 1998) and Pseudo code for SMO (Platt, 1998)
- ❄ Motivation: bad performance of  $\text{svm}^{\text{light}}$  on complex cases  $\Rightarrow$  development of `bsvm`
- ❄ Drawbacks of `bsvm`: very complex, uses third-party optimization solvers  $\Rightarrow$  simple, usable code: `libsvm`
- ❄ Meanwhile: several extensions and optimizations  $\Rightarrow$  very competitive
- ❄ Future development: increase usability (“black box”)

# Features

---

❄  $C$  and  $\nu$  classification

❄ multiclass-classification

❄  $\epsilon$  and  $\nu$  regression

❄ novelty detection

❄ unbalanced data

❄ internal sparse data format

❄ many interfaces (C, Python, Java, Perl, R, MATLAB, ...)

# Performance

---

- ❄ libsvm won several competitions, e.g.:
  - ❖ EUNITE 2001 (competition on electricity load prediction)
  - ❖ IJCNN Challenge 2001 (two of three competitions)
- ❄ very good (fast) quadratic solver (decomposition method), high standard in numerical methods.

# The R interface

---

❄ libsvm is a C++ library with standard C interface


❄ R interface in package e1071 on CRAN  
(<http://cran.R-project.org/>)

❄ main functions:

- ❖ training (returning an 'svm' object)
- ❖ predicting
- ❖ tuning (for hyperparameter selection)

# Training

$$y = \log(x) + \cos(x) + \varepsilon$$



# Training

---

```
R> library(e1071)
R> f <- function(x) log(x) + cos(x)
R> x <- seq(0.1, 5, by = 0.05)
R> y <- f(x) + rnorm(x, sd = 0.2)
R> print(svmmodel <- svm(x, y))
```

Call:

```
svm.default(x = x, y = y)
```

Parameters:

SVM-Type: eps-regression

SVM-Kernel: radial


cost: 1

gamma: 1

epsilon: 0.1

# Training

$$y = \log(x) + \cos(x) + \varepsilon$$



# Parameter Tuning

---

```
R> tunemodel <- tune.svm(x, y, gamma = 2^(-4:0), cost = 2^(-2:2))  
R> tunemodel
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

gamma	cost
-------	------


2	16
---	----

- best performance: 0.05016086

```
R> plot(tunemodel)
```




# Parameter Tuning



# Parameter Tuning

$$y = \log(x) + \cos(x) + \varepsilon$$



# libsvm under test

---

- ❄ No extensive benchmark study available comparing SVMs to many methods on many data sets
- ❄ Often only compared to other machine learning techniques (e.g., neural networks)
- ❄ How do SVMs compare to 'traditional' methods (e.g., linear models)?
- ❄ R offers a wide variety of methods, all accessible with a similar interface  $\Rightarrow$  easy to use in a benchmark study

# Methods

---

9 regression methods:

**Machine Learning Techniques:** Neural Networks, Trees

**Resample & Combine Methods:** Bagged Trees, Random Forests, MART

**Spline Models:** MARS, BRUTO, Projection Pursuit

**‘Traditional’ Methods:** Linear Models

# Data sets

---

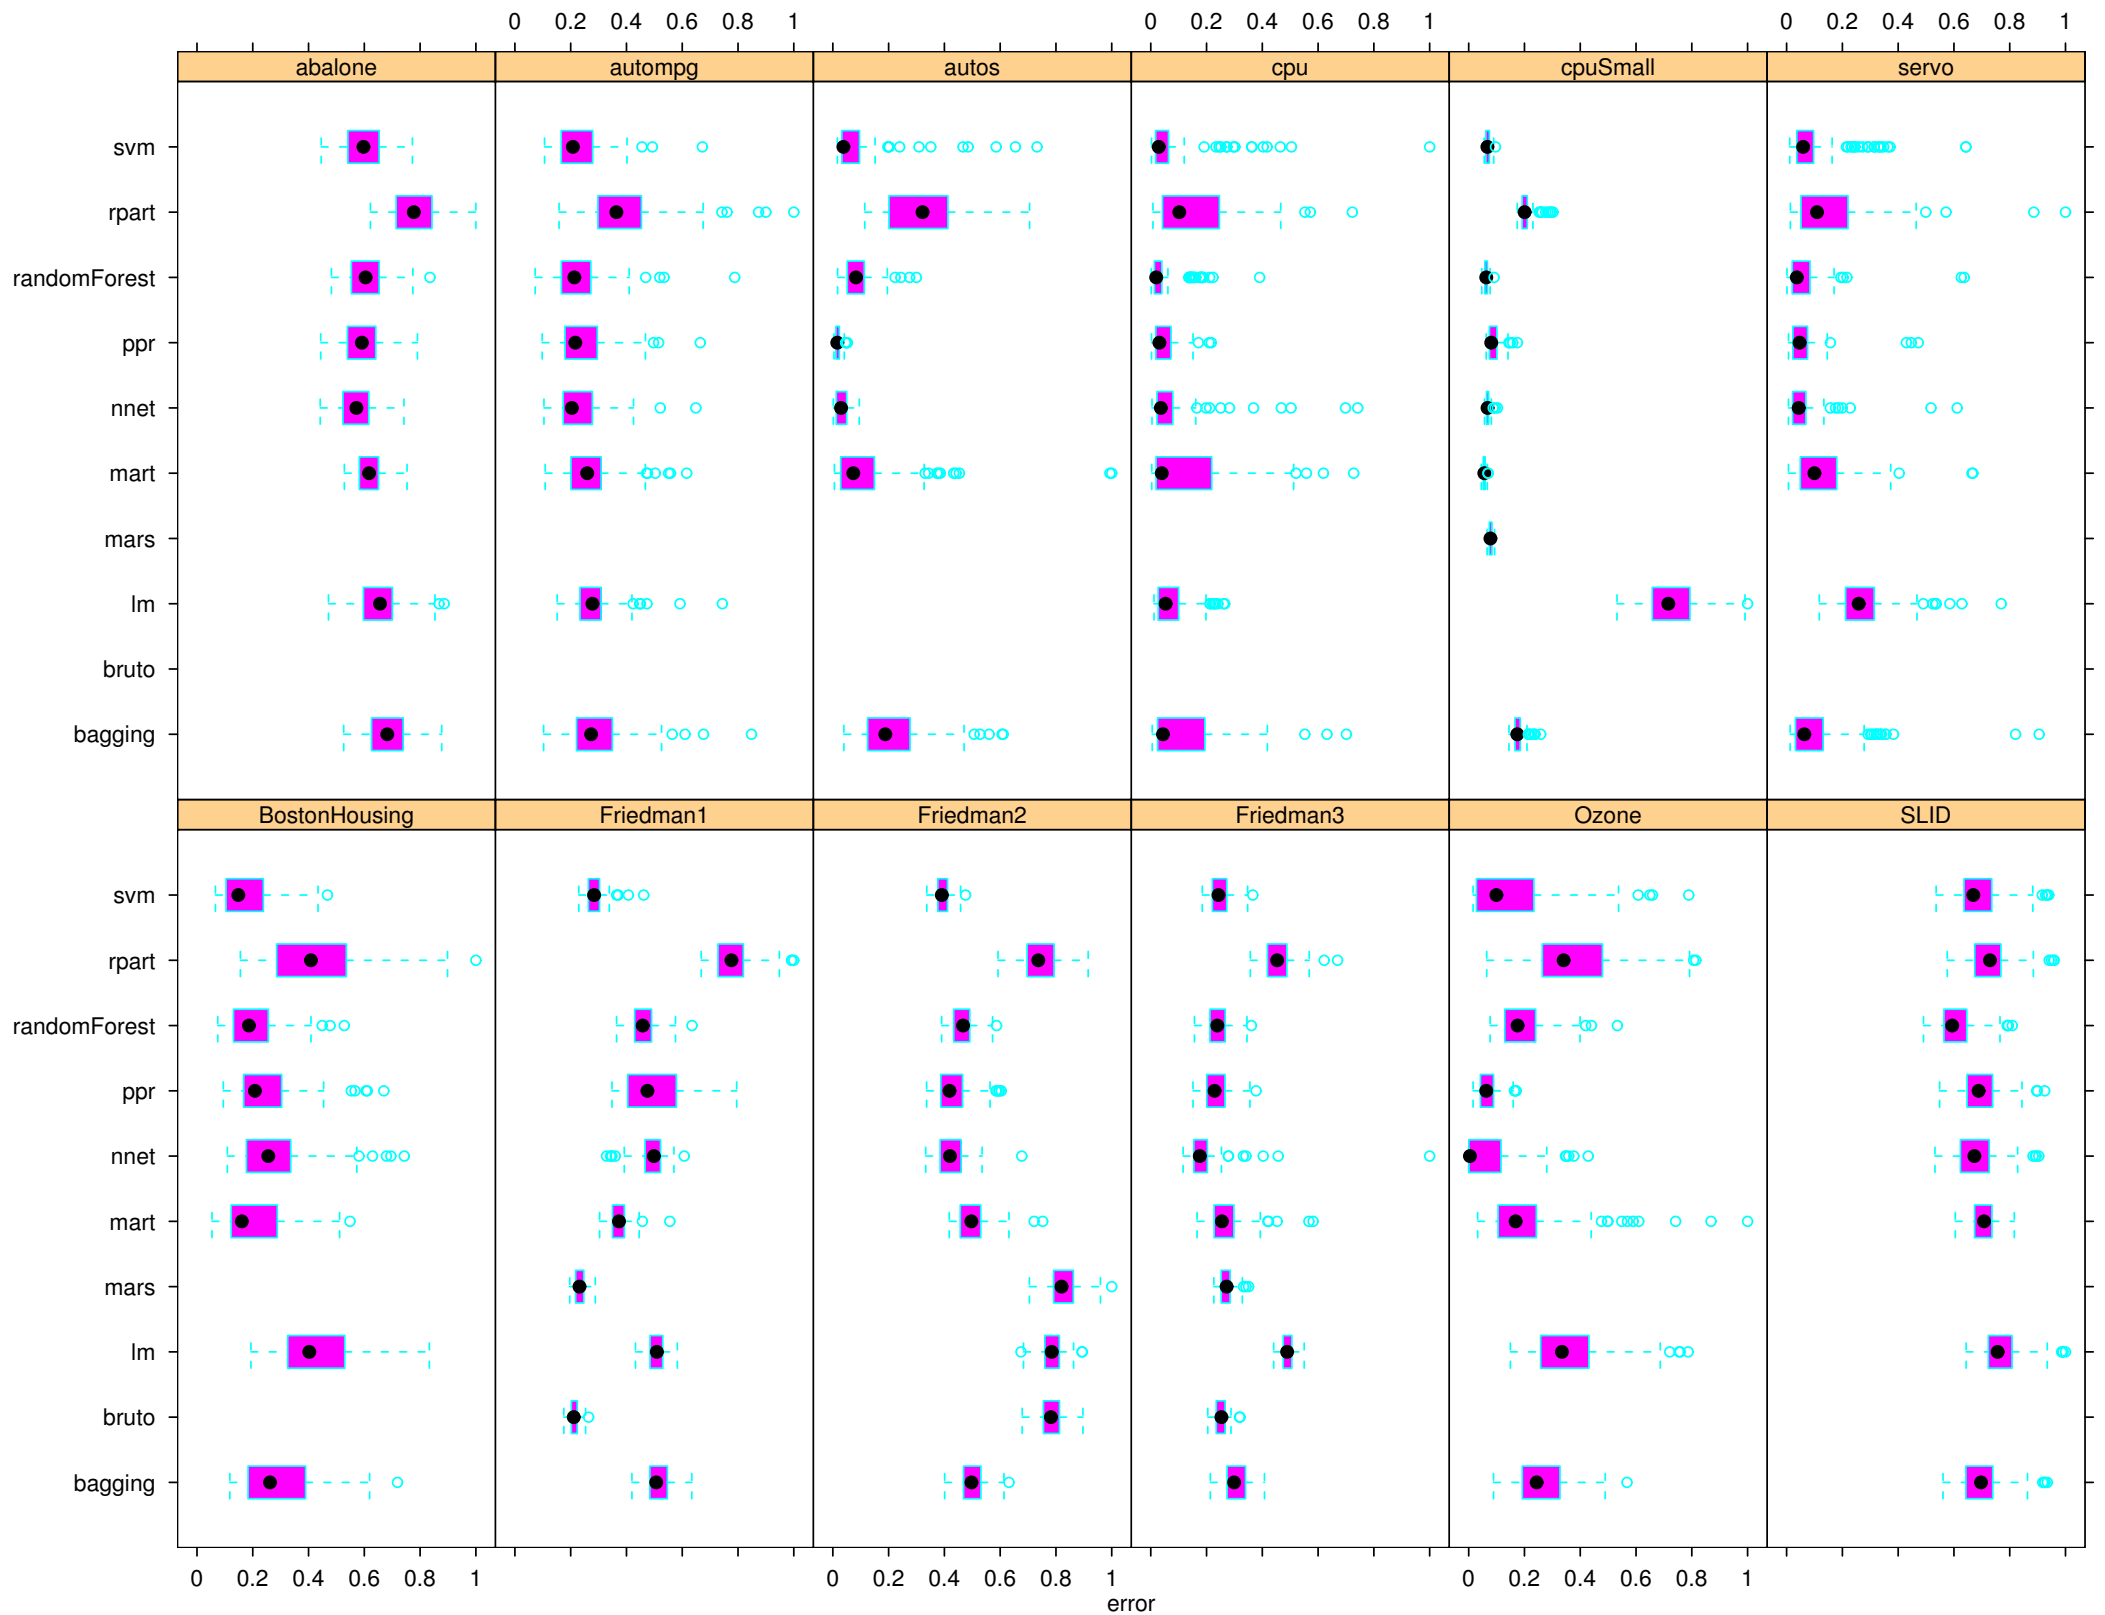
Most data sets were taken from the UCI machine learning data base:

- ❄ 12 data sets
- ❄ artificial and real data
- ❄ small (167) to large (8192) data sets
- ❄ Mix of binary, categorical, and metric input variables


# Simulation Setup

---

- ❄ Generation of 100 training and 100 test sets  
(for real data, using 10 times 10-fold cross validation)
- ❄ Scaling of data
- ❄ Tuning on  $\frac{2}{3}$  of training set, other  $\frac{1}{3}$  used for validation
- ❄ Training on complete training set
- ❄ Performance measure computed on test set  
(Mean Squared Error)



# Results





# Results

---

- ❄ Generally good performance of SVMs (almost always ranked in top 3)
- ❄ However, only ranked first on two data sets
- ❄ Average rankings for dispersion
- ❄ Good performances by neural networks and random forests

# The EUNITE competition

---

World-wide competition on electricity load prediction (Eastern Slovakia)

**Data:** loads (every 30min) and temperatures for 2 years (1997 and 1998)


**Task:** predict maximum load for a whole month (January 1999)

**Performance:** Mean Average Prediction Error

$$\text{MAPE\%} = 100 \times \sum_{t=1}^T \left| \frac{y_t - \hat{y}_t}{y_t} \right|$$


# The EUNITE competition

Load Maxima for 1997 and 1998



# The EUNITE competition

Observed Data in January 1999



# The EUNITE competition

---

## Data preparation:

- ❄ Maximum loads of seven past days
- ❄ Seven binary attributes for day of week
- ❄ One binary attribute for indicating a holiday
- ❄ One attribute for daily average temperature

After some experiments, some information has been discarded:

- both holiday and temperature
- data in summer time (April–September)

# The EUNITE competition

---

## Training:

1. Use January, 1998, as test set, rest as training set
2. Hyperparameters:
  - ❄ use RBF kernel ( $K(x_i, x_j) = \exp -\gamma \|x_i - x_j\|^2$ )
  - ❄ use  $\epsilon = 0.5$  (default)
  - ❄ tune  $C$  and  $\gamma$  using 5-fold cross-validation

# The EUNITE competition

---


## **Predicting:**

1. Start with loads of last seven days of December, 1998
2. Remove first value, add predicted value for next day
3. Repeat until all values are predicted

# The EUNITE competition


---

True and Predicted Data for January 1999





# The EUNITE competition



# Resources

---



CRAN:

<http://cran.R-project.org/>



libsvm:

<http://www.csie.ntu.edu.tw/~cjlin/libsvm/>



EUNITE:

<http://neuron.tuke.sk/competition/>



Benchmarks:

<http://www.wu-wien.ac.at/am/Download/report78.pdf>



This talk & data sets:

<http://www.ci.tuwien.ac.at/~meyer/svm/>